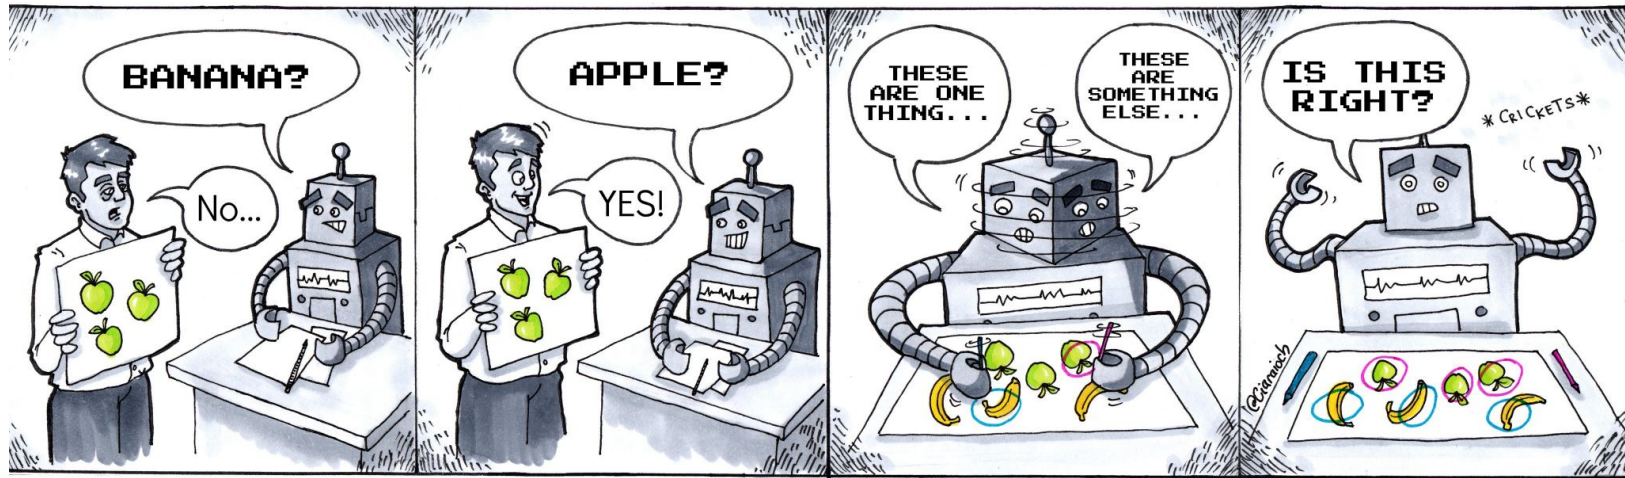


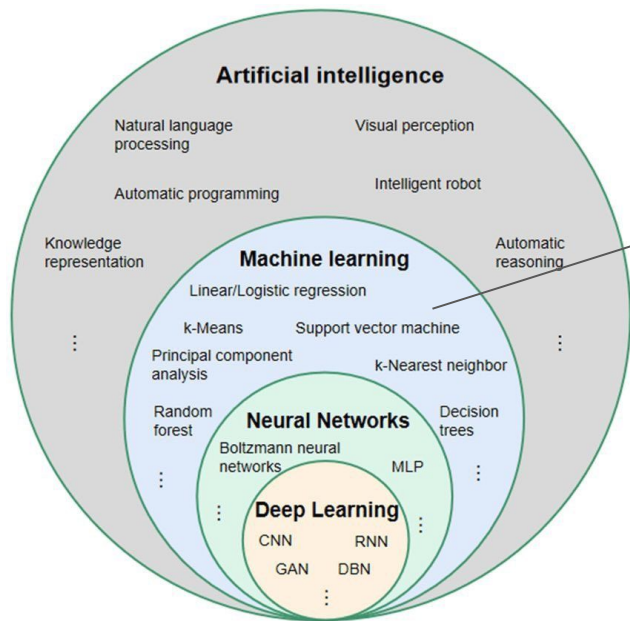
Machines that Learn



Supervised Learning

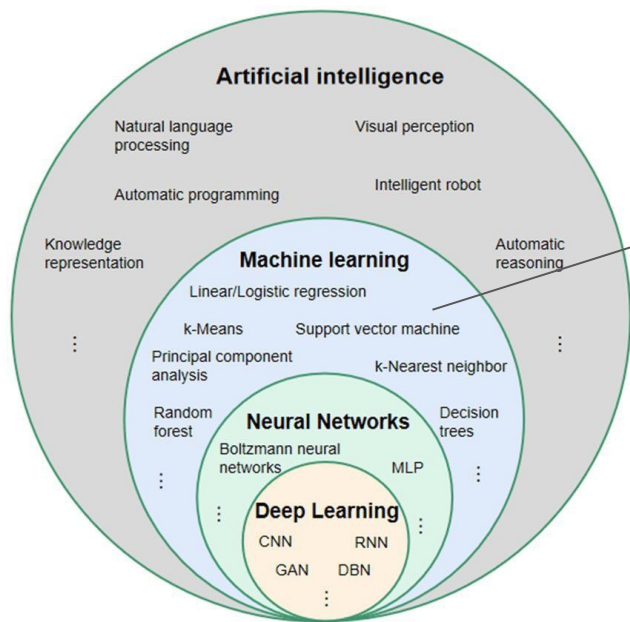
Unsupervised Learning

Recap



Supervised Learning
Unsupervised Learning
Reinforcement Learning

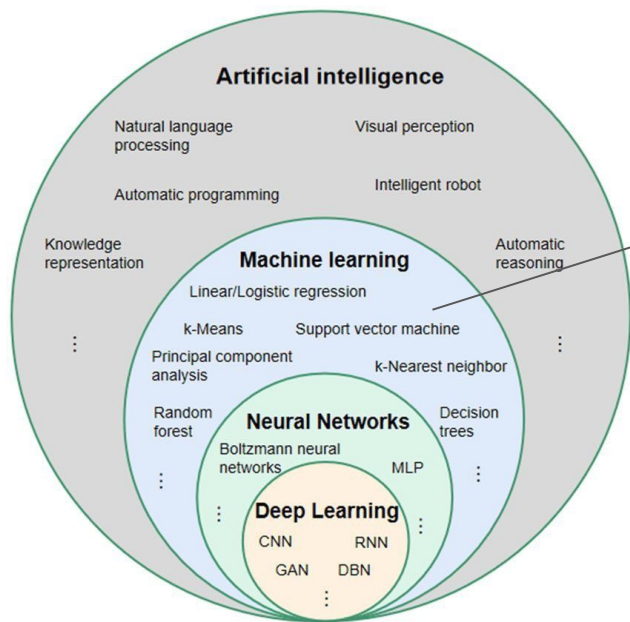
Recap



Supervised Learning
Self-Supervised Learning
Few-shot Learning
Zero-shot Learning
Transfer Learning
Unsupervised Learning
Reinforcement Learning

Federated Learning

Recap



Supervised Learning

Self-Supervised Learning

Few-shot Learning

Zero-shot Learning

Transfer Learning

Unsupervised Learning

Reinforcement Learning

Federated Learning

Recap

Supervised Learning

$$f(x) = y$$

Recap

Supervised Learning

$$f(x) = y$$



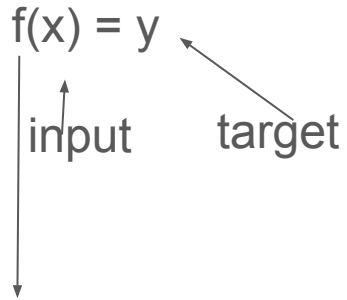
input



target

Recap

Supervised Learning



Machine learning algorithm i.e logistic classification, logistic regression, naive bayes, etc...

Recap

Supervised Learning

$$f(x) = y$$

input

target

Tries to match the prediction to the targets

We have the targets

Machine learning algorithm i.e logistic classification, logistic regression, naive bayes, etc...

Recap

Supervised Learning

$$f(x) = y$$

Getting this

P.s we don't always have labels readily available

Making sure the data is good

Finding the most suitable algorithm and training it

Recap

Supervised Learning

$$f(x) = y$$

```
graph LR; A["f(x) = y"] --> B["Getting this"]; A --> C["Evaluate if the result makes sense and is relevant for the given use-case"]; A --> D["Making sure the data is good"]; A --> E["Finding the most suitable algorithm and training it"];
```

The diagram illustrates the supervised learning process. It starts with a central box containing the equation $f(x) = y$. From this box, four arrows point outwards to different stages of the process: a horizontal arrow to the right labeled 'Getting this', a curved arrow pointing down and to the right labeled 'Making sure the data is good', a vertical arrow pointing down labeled 'Finding the most suitable algorithm and training it', and a long horizontal arrow pointing to the right labeled 'Evaluate if the result makes sense and is relevant for the given use-case'.

Getting this

Evaluate if the result makes sense and is relevant for the given use-case

Making sure the data is good

Finding the most suitable algorithm and training it

Supervised Learning

This class will be about

- Being empiricists
- (more) Exploratory Data Analysis (EDA)
- Feature Engineering (more data wrangling)
- Training Models
- Evaluation

Supervised Learning

If machines can learn by themselves, just give it data.
Shouldn't it be easy ?

Supervised Learning

If machines can learn by themselves, just give it data.
Shouldn't it be easy ?



Supervised Learning

How do we separate good from garbage ?

Supervised Learning

How do we separate good from garbage ?
How do we know if it smells ?



<https://ideogram.ai/>

Supervised Learning

How do we separate good from garbage ?

How do we know if it smells ?

Exploratory data analysis, explore correlation and quickly build baseline models

Supervised Learning

How do we separate good from garbage ?
How do we know if it smells ?

Things to aim for when working on a dataset:

- Complete
- Consistent
- Relevant
- Representative
- Accurate

Supervised Learning

How do we separate good from garbage ?

How do we know if it smells ?

Things to aim for when working on a dataset:

- **Complete** - Ensure all necessary data points are included and address missing data appropriately.
- **Consistent** - Maintain uniform formats and values throughout the dataset to avoid contradictions.
- **Relevant** - Include only data directly aligned with your analysis goals or application needs
- **Representative** - Ensure the dataset accurately reflects the population or phenomenon being studied.
- **Accurate** - Verify data correctness to minimize errors or misleading insights

Supervised Learning

How do we separate good from garbage ?

How do we know if it smells ?

Things to aim for when working on a dataset:

- Complete
- Consistent
- Relevant
- Representative
- Accurate

Red flags

- Too many empties or NaN
- Outliers
- Biases
- Noisy Data
- Weird things

Supervised Learning

Data → Features

Supervised Learning

Feature Engineering

- Categorical Encoding
- Handle Missing Values
- Feature Scaling
- Creating new features

Data → Features

Supervised Learning

Feature Engineering

- Categorical Encoding
 - one-hot encoding
 - multi-hot encoding

Data → Features

Supervised Learning

Feature Engineering

Data → Features

- Categorical Encoding
 - one-hot encoding
 - multi-hot encoding

Some ways to do it:

- [df.get_dummies](#)
- [sklearn one hot encoder](#)
- [sklearn multi hot encoding](#)

Supervised Learning

Feature Engineering

- Handle Missing Values

Data → Features

Supervised Learning

Feature Engineering

Data → Features

- Handle Missing Values
 - Drop
 - Imputation (Mean, Median, Mode)
 - Forward Fill and Backward Fill
 - Random Values

[Top 4 Techniques for Handling Missing Values in Machine Learning](#)

Supervised Learning

Feature Engineering

- Feature Scaling
 - Normalization
 - Standardization

Data → Features

Supervised Learning

Feature Engineering

Data → Features

- Creating new features based on existing data

Supervised Learning

Cool ?

Supervised Learning



Supervised Learning

Data understood.
Features created.
Now it's time to train.



Supervised Learning

Linear Regression
Logistic Regression



Supervised Learning

Linear Regression

Logistic Regression

Naive-Bayes

KNN (k-nearest neighbors)

SVM (Support Vector Machines)



Supervised Learning

Naive-Bayes

Naive Bayes is a **probabilistic** machine learning algorithm based on **Bayes' Theorem**, which calculates the **probability** of a class given a set of features. It assumes that all features are independent (the "naive" assumption), which simplifies computation but may not hold true in all real-world cases. Despite this limitation, it is highly effective for **text classification** and **spam filtering** because it works well with **high-dimensional data** and **provides fast predictions**.

Reference: [Sklearn Naive Bayes](#)

Supervised Learning

Naive-Bayes

Naive Bayes is a **probabilistic** machine learning algorithm based on **Bayes' Theorem**, which calculates the **probability** of a class given a set of features. It assumes that all features are independent (the "naive" assumption), which simplifies computation but may not hold true in all real-world cases. Despite this limitation, it is highly effective for **text classification** and **spam filtering** because it works well with **high-dimensional data** and **provides fast predictions**.

Reference: [Sklearn Naive Bayes](#)

Challenge: Try to implement this algorithm from scratch

Supervised Learning

K-Nearest Neighbors (KNN)

KNN is a simple, instance-based algorithm that **predicts the class** of a data point by **considering the classes of its k closest neighbors in the feature space**. The "closeness" is typically determined using distance metrics like **Euclidean distance**. KNN is intuitive and non-parametric, meaning it makes no assumptions about the underlying data distribution, but it can **become computationally expensive as the dataset grows**.


Reference: [Sklearn KNN](#)

Supervised Learning

K-Nearest Neighbors (KNN)

KNN is a simple, instance-based algorithm that **considering the classes of its k closest neighbors**. "closeness" is typically determined using distance metrics. KNN is intuitive and non-parametric, meaning it does not assume an underlying data distribution, but it can **become inefficient** as the **dataset grows**.

Reference: [Sklearn KNN](#)



Supervised Learning

Support Vector Machines (SVM)

SVM is a powerful algorithm used for **classification** and **regression** tasks. It works by finding a **hyperplane that best separates data points** of different classes in a **high-dimensional space**. SVM focuses on maximizing the margin between the classes, making it robust to outliers and well-suited for both linear and non-linear data (with the help of kernel functions). **It is especially effective for datasets with clear margins of separation.**

Reference: [Sklearn SVM](#)

Let's do some feature engineering

More EDA
Feature Engineering
More training



Let's do some feature engineering

More EDA
Feature Engineering
More training



Supervised Learning

Remember when I said we need to be empiricists ?

And act like scientists.

We need to create hypothesis and properly evaluate them.



Supervised Learning

Data = Train + Testing

Testing = Validation + Testing

Data = Train + Validation + Testing



Supervised Learning

Data = Train + Testing

Testing = **Validation** + Testing

Data = Train + **Validation** + Testing

Validation - used during development to validate hyperparameters (like max_iter of logistic regression)

* more on this on the next class



Supervised Learning

Data = Train + Testing

Testing = Validation + Testing

Data = Train + Validation + Testing

Sidenote:

In real world use-cases data is not stagnant.

Any ML models need to be constantly re-evaluated
and more often than not re-trained.



Supervised Learning

Hypothesis

We can learn a reasonable enough valid correlation between X and Y

Testing

How close did we get ?
Is it close enough ?



Supervised Learning

Testing results are only considered correct if there are no biases* during training or evaluation - i.e no **data leakage**

*data leakage is just one possible bias, we also can have selection bias, sampling bias among others



[Understanding Bias in Machine Learning Models - Arize AI](#)

Supervised Learning

Testing results are only considered correct if there are no biases during training or evaluation - i.e no **data leakage**

This means that testing data, *ideally* is locked away.



Supervised Learning

Testing results are only considered correct if there are no biases during training or evaluation - i.e no **data leakage**

This means that testing data, *ideally* is locked away.

Which means that certain **transformations** must be done **only after splitting the data** e.g Normalization, Standardization, TF-IDF,...



Supervised Learning

Hypothesis

We can learn a reasonable enough valid correlation between X and Y

Testing

How close did we get ?

Is it close enough ?



Supervised Learning

Evaluation is looking closely at the right metrics.

Let's go back to titanic and store sales.



Supervised Learning

Hypothesis

We can learn a reasonable enough valid correlation between X and Y

Testing

How close did we get ?

Is it close enough ?



Supervised Learning


Building blocks

- [Pipeline — scikit-learn 1.5.2 documentation](#)
- [sklearn.datasets — scikit-learn 1.5.2 documentation](#)
- [sklearn.preprocessing — scikit-learn 1.5.2 documentation](#)




Machines that Learn


More on the next class



Supervised Learning



Supervised Learning



Supervised Learning

Hyperparameter Search
Cross-Validation